



ara218 / Ara218



- <> Code
- Issues
- Pull requests
- Projects
- Wiki
- Security
- Insights
- Settings

main



Ara218 / HomeWork / HW_6.ipynb

ara218 homework 6 done

History

1 contributor

1.74 MB

HW 6

This assignment covers all fundamental concepts required for completing a project

DO NOT ERASE MARKDOWN CELLS AND INSTRUCTIONS IN YOUR HW submission

- **Q** - QUESTION
- **A** - Where to input your answer

Instructions

Keep the following in mind for all notebooks you develop:

- Structure your notebook.
- Use headings with meaningful levels in Markdown cells, and explain the questions each piece of code is to answer or the reason it is there.
- Make sure your notebook can always be rerun from top to bottom.
- Please start working on this assignment as soon as possible. If you are a beginner in Python this might take a long time. One of the objectives of this assignment is to help you learn python and scikit-learn package.
- See [README.md](#) for homework submission instructions

Related Tutorials

Refreshers

- [Intro to Machine Learning w scikit-learn](#)
- [A tutorial on statistical-learning for scientific data processing](#)

Classification Approaches

- [Logistic Regression with Sklearn](#)
- [KNN with sklearn](#)
- [Support Vector machine example](#)
- [SVC](#)
- [Bagging Classifier](#)
- [Gradient Boosting Classifier](#)

Modeling

- [Cross-validation](#)
- [Plot Confusion Matrix with Sklearn](#)
- [Confusion Matrix Display](#)

Import all required library

In [4]:

```
import pandas as pd
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import numpy as np
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, MaxAbsScaler
import json
import lightgbm as lgbm
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
import seaborn as sns
from imblearn.over_sampling import RandomOverSampler
from sklearn.ensemble import RandomForestClassifier
```

Data Processing

Q1 Get training data from the dataframe

1. Load HW6_data.csv from data folder into data frame
2. Print the head of the dataframe
3. Print the shape of the dataframe
4. Print the description of the dataframe
5. Check if the dataset has NULL values. (Show number of NULL values per column)
6. Assign Cover_Type values to Y
7. Assign rest of the column values to X

A1 Fill the cell blocks below, Create new cell as per your necessary

In [7]:

```
#You can create or remove cells as per your need
df = pd.read_csv('C:\\Users\\alsae\\Desktop\\fake\\2024Spring\\data\\HW6_data.csv')
print(df.head(5))
print(df.shape)
print(df.describe())
print(df.isnull().sum())
```

	Elevation	Aspect	Slope	Horizontal_Distance_To_Hydrology	\
0	3080.0	137	18.0		166
1	2758.0	19	8.0		551
2	2779.0	86	9.0		43
3	2811.0	296	0.0		287
4	2956.0	314	26.0		71

	Vertical_Distance_To_Hydrology	Horizontal_Distance_To_Roadways	\
0		1	1009
1		49	1766
2		-10	3889
3		4	788
4		22	2010

```

Hillshade_9am Hillshade_Noon Hillshade_3pm \
0          250.0          198          166
1          225.0          231          124
2          155.0          204          123
3          191.0          226          113
4          230.0          200          99

Horizontal_Distance_To_Fire_Points ... Soil_Type32 Soil_Type33 \
0          3635.0 ...          0          0
1          1648.0 ...          0          0
2          364.0 ...          0          0
3          144.0 ...          0          0
4          743.0 ...          0          0

Soil_Type34 Soil_Type35 Soil_Type36 Soil_Type37 Soil_Type38 \
0          0          0.0          0          0          0
1          0          0.0          0          0          0
2          0          0.0          0          0          1
3          0          0.0          0          0          0
4          0          0.0          0          0          1

Soil_Type39 Soil_Type40 Cover_Type
0          0          0.0          1
1          0          0.0          2
2          0          0.0          2
3          0          0.0          2
4          0          NaN          2

[5 rows x 55 columns]
(80000, 55)

Elevation Aspect Slope \
count 79433.000000 80000.000000 79346.000000
mean 2981.436531 151.634175 15.092494
std 287.979705 109.945631 8.528153
min 1813.000000 -29.000000 -3.000000
25% 2762.000000 60.000000 9.000000
50% 2967.000000 122.000000 14.000000
75% 3217.000000 246.000000 20.000000
max 4271.000000 400.000000 61.000000

Horizontal_Distance_To_Hydrology Vertical_Distance_To_Hydrology \
count 80000.000000 80000.000000
mean 271.564212 51.510737
std 227.532197 68.091489
min -43.000000 -276.000000
25% 111.000000 4.000000
50% 212.000000 31.000000
75% 361.000000 78.000000
max 1544.000000 562.000000

Horizontal_Distance_To_Roadways Hillshade_9am Hillshade_Noon \
count 80000.000000 79200.000000 80000.000000
mean 1770.080712 211.786818 221.069125
std 1318.661060 30.822278 22.191030
min -238.000000 10.000000 69.000000
25% 821.000000 198.000000 210.000000
50% 1440.000000 218.000000 224.000000
75% 2366.000000 233.000000 237.000000
max 7604.000000 293.000000 264.000000

```

	Hillshade_3pm	Horizontal_Distance_To_Fire_Points	...	Soil_Type32	\
count	80000.000000	78870.000000	...	80000.000000	
mean	140.711750	1578.058615	...	0.038150	
std	43.859689	1125.780446	...	0.191559	
min	-48.000000	-218.000000	...	0.000000	
25%	115.000000	782.000000	...	0.000000	
50%	142.000000	1362.000000	...	0.000000	
75%	169.000000	2082.000000	...	0.000000	
max	268.000000	8011.000000	...	1.000000	

	Soil_Type33	Soil_Type34	Soil_Type35	Soil_Type36	Soil_Type37	\
count	80000.000000	80000.000000	79720.000000	80000.000000	80000.000000	
mean	0.037687	0.011838	0.015429	0.010812	0.012538	
std	0.190441	0.108155	0.123252	0.103420	0.111268	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	1.000000	

	Soil_Type38	Soil_Type39	Soil_Type40	Cover_Type
count	80000.000000	80000.000000	75000.000000	80000.000000
mean	0.040325	0.039163	0.030707	1.770725
std	0.196722	0.193983	0.172523	0.892577
min	0.000000	0.000000	0.000000	1.000000
25%	0.000000	0.000000	0.000000	1.000000
50%	0.000000	0.000000	0.000000	2.000000
75%	0.000000	0.000000	0.000000	2.000000
max	1.000000	1.000000	1.000000	7.000000

[8 rows x 55 columns]

Elevation	567
Aspect	0
Slope	654
Horizontal_Distance_To_Hydrology	0
Vertical_Distance_To_Hydrology	0
Horizontal_Distance_To_Roadways	0
Hillshade_9am	800
Hillshade_Noon	0
Hillshade_3pm	0
Horizontal_Distance_To_Fire_Points	1130
Wilderness_Area1	0
Wilderness_Area2	0
Wilderness_Area3	0
Wilderness_Area4	0
Soil_Type1	0
Soil_Type2	0
Soil_Type3	0
Soil_Type4	0
Soil_Type5	0
Soil_Type6	0
Soil_Type7	0
Soil_Type8	0
Soil_Type9	0
Soil_Type10	0
Soil_Type11	0
Soil_Type12	0
Soil_Type13	0
Soil_Type14	0

```

Soil_Type15      0
Soil_Type16      0
Soil_Type17      0
Soil_Type18      0
Soil_Type19      0
Soil_Type20      0
Soil_Type21      0
Soil_Type22      0
Soil_Type23      0
Soil_Type24      0
Soil_Type25      0
Soil_Type26      0
Soil_Type27      0
Soil_Type28      0
Soil_Type29      0
Soil_Type30      0
Soil_Type31      0
Soil_Type32      0
Soil_Type33      0
Soil_Type34      0
Soil_Type35      280
Soil_Type36      0
Soil_Type37      0
Soil_Type38      0
Soil_Type39      0
Soil_Type40      5000
Cover_Type       0
dtype: int64

```

In [8]:

```

X=df.drop(columns=['Cover_Type'])
Y=df['Cover_Type']

```

Q2: Open-Ended Questions: Observe the range of all feature values and statistical information from the dataframe description above.

1. If the dataset has NULL values, Give proper justification about the methods you will use to replace NULL values for specific columns.
2. Do you think in our dataset normalization is required? -- Give proper justification based on your opinion.
3. What type of normalization/Scaling technique you whould recommend for our dataset?

A2

Answer 1:

Answer 2:

Answer 3:

Q3:

1. Replace the null values with the best possible methods from your above observation
2. Use the above mentioned normalization technique on our HW_6 dataset.
3. Transform the X dataframe using chosen normalization technique

5. transform the X dataframe using choosen normalization technique.

Note: Make sure the scaled X has all column name same as X dataframe

A3 Fill the cell blocks below, Create new cell as per your necessary

```
In [9]: for col in X.columns:
        if X[col].dtype == 'object':
            X[col] = X[col].fillna(X[col].mode()[0])
        else:
            X[col] = X[col].fillna(X[col].mean())
```

```
In [10]: from sklearn.preprocessing import MinMaxScaler

# Initialize MinMaxScaler
scaler = MinMaxScaler()


# Fit and transform the data
Scaled_X = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
```

```
In [11]: Scaled_X
```

```
Out[11]:
```

	Elevation	Aspect	Slope	Horizontal_Distance_To_Hydrology	Vertical_Distanc
0	0.515460	0.386946	0.328125		0.131695
1	0.384459	0.111888	0.171875		0.374291
2	0.393002	0.268065	0.187500		0.054190
3	0.406021	0.757576	0.046875		0.207940
4	0.465012	0.799534	0.453125		0.071834
...
79995	0.288853	0.188811	0.171875		0.080655
79996	0.370627	0.100233	0.218750		0.261500
79997	0.557771	0.552448	0.109375		0.335224
79998	0.635883	0.081585	0.359375		0.299937
79999	0.567535	0.606061	0.234375		0.101449

80000 rows × 54 columns



Q4:

1. Check again and show if there is any null values left in our Scaled_X .
2. Print all unique values/ different class id from the Y data .

A4 Fill the cell blocks below, Create new cell as per your necessary

In [12]:

```
print("NULL values in Scaled_X:")
print(Scaled_X.isnull().sum())
```

```
NULL values in Scaled_X:
Elevation                                0
Aspect                                  0
Slope                                  0
Horizontal_Distance_To_Hydrology        0
Vertical_Distance_To_Hydrology          0
Horizontal_Distance_To_Roadways         0
Hillshade_9am                          0
Hillshade_Noon                         0
Hillshade_3pm                         0
Horizontal_Distance_To_Fire_Points      0
Wilderness_Area1                      0
Wilderness_Area2                      0
Wilderness_Area3                      0
Wilderness_Area4                      0
Soil_Type1                            0
Soil_Type2                            0
Soil_Type3                            0
Soil_Type4                            0
Soil_Type5                            0
Soil_Type6                            0
Soil_Type7                            0
Soil_Type8                            0
Soil_Type9                            0
Soil_Type10                           0
Soil_Type11                           0
Soil_Type12                           0
Soil_Type13                           0
Soil_Type14                           0
Soil_Type15                           0
Soil_Type16                           0
Soil_Type17                           0
Soil_Type18                           0
Soil_Type19                           0
Soil_Type20                           0
Soil_Type21                           0
Soil_Type22                           0
Soil_Type23                           0
Soil_Type24                           0
Soil_Type25                           0
Soil_Type26                           0
Soil_Type27                           0
Soil_Type28                           0
Soil_Type29                           0
Soil_Type30                           0
Soil_Type31                           0
Soil_Type32                           0
Soil_Type33                           0
Soil_Type34                           0
Soil_Type35                           0
Soil_Type36                           0
Soil_Type37                           0
Soil_Type38                           0
```



```
Soil_Type39      0
Soil_Type40      0
dtype: int64
```

```
In [13]: # Print unique values from Y
print("\nUnique values in Y:")
print(Y.unique())
```

```
Unique values in Y:
[1 2 3 7 6 4]
```

Part 1: Use a subset of whole data(N=20000) for Data Visualization

Data Subset Creation

1. First we are Selecting $N=20000$ random rows from our original dataset which is `df` and create a new subset of data.
2. Using the below **rndperm** and selecting first N index from the `Scaled_X` and `Y`

```
In [16]: np.random.seed(42)
rndperm = np.random.permutation(df.shape[0])
N = 20000
data_subset_x = Scaled_X.loc[rndperm[:N],:].copy()
data_subset_y = Y.loc[rndperm[:N]].copy()
```

Q5:

1. Use PCA and reduce the dimension of the **data_subset_x** into 3 .
2. Store the PCA result into `pca_result` variable
3. Add the results from the PCA into the **data_subset_x** as new columns. (Choose any meaningful names for the columns)

A5 Fill the below cells. Use extra cells as per your necessary

```
In [17]: #You can create or remove cells as per your need
from sklearn.decomposition import PCA

pca = PCA(n_components=3)
pca_result = pca.fit_transform(data_subset_x)
```

```
In [18]: data_subset_x['PCA_Component_1'] = pca_result[:, 0]
data_subset_x['PCA_Component_2'] = pca_result[:, 1]
data_subset_x['PCA_Component_3'] = pca_result[:, 2]
```

Q6:

1. Use TSNE and reduce the dimension of the **data_subset_x** into 2 .

2. Store the TSNE result into `tsne_results` variable
3. Add the results from the T-SNE into the **data_subset_x** as new columns. (Choose any meaningful names for the columns)

Note:

1. You can use `from sklearn.manifold import TSNE` for TSNE initialization.
2. Give value of `n_components` as per the question.
3. Also use other parameters while TSNE initialization as, `verbose=1`, `perplexity=40`, `n_iter=300`

A6 Fill the below cells. Use extra cells as per your necessary

In [19]:

```
#You can create or remove cells as per your need
from sklearn.manifold import TSNE

tsne = TSNE(n_components=2, verbose=1, perplexity=40, n_iter=300)
tsne_results = tsne.fit_transform(data_subset_x)
```

```
[t-SNE] Computing 121 nearest neighbors...
[t-SNE] Indexed 20000 samples in 0.006s...
[t-SNE] Computed neighbors for 20000 samples in 1.386s...
[t-SNE] Computed conditional probabilities for sample 1000 / 20000
[t-SNE] Computed conditional probabilities for sample 2000 / 20000
[t-SNE] Computed conditional probabilities for sample 3000 / 20000
[t-SNE] Computed conditional probabilities for sample 4000 / 20000
[t-SNE] Computed conditional probabilities for sample 5000 / 20000
[t-SNE] Computed conditional probabilities for sample 6000 / 20000
[t-SNE] Computed conditional probabilities for sample 7000 / 20000
[t-SNE] Computed conditional probabilities for sample 8000 / 20000
[t-SNE] Computed conditional probabilities for sample 9000 / 20000
[t-SNE] Computed conditional probabilities for sample 10000 / 20000
[t-SNE] Computed conditional probabilities for sample 11000 / 20000
[t-SNE] Computed conditional probabilities for sample 12000 / 20000
[t-SNE] Computed conditional probabilities for sample 13000 / 20000
[t-SNE] Computed conditional probabilities for sample 14000 / 20000
[t-SNE] Computed conditional probabilities for sample 15000 / 20000
[t-SNE] Computed conditional probabilities for sample 16000 / 20000
[t-SNE] Computed conditional probabilities for sample 17000 / 20000
[t-SNE] Computed conditional probabilities for sample 18000 / 20000
[t-SNE] Computed conditional probabilities for sample 19000 / 20000
[t-SNE] Computed conditional probabilities for sample 20000 / 20000
[t-SNE] Mean sigma: 0.184852
[t-SNE] KL divergence after 250 iterations with early exaggeration: 80.914803
[t-SNE] KL divergence after 300 iterations: 2.991430
```

In [20]:

```
data_subset_x['tSNE_Component_1'] = tsne_results[:, 0]
data_subset_x['tSNE_Component_2'] = tsne_results[:, 1]
```

Q7:

1. Create a new dataframe with name `df_plot`
2. This dataframe will merge everything from **data_subset_x** and **data_subset_y**
3. We need to give a name for the `data_subset_y` column. Use `Cover_Type` as the

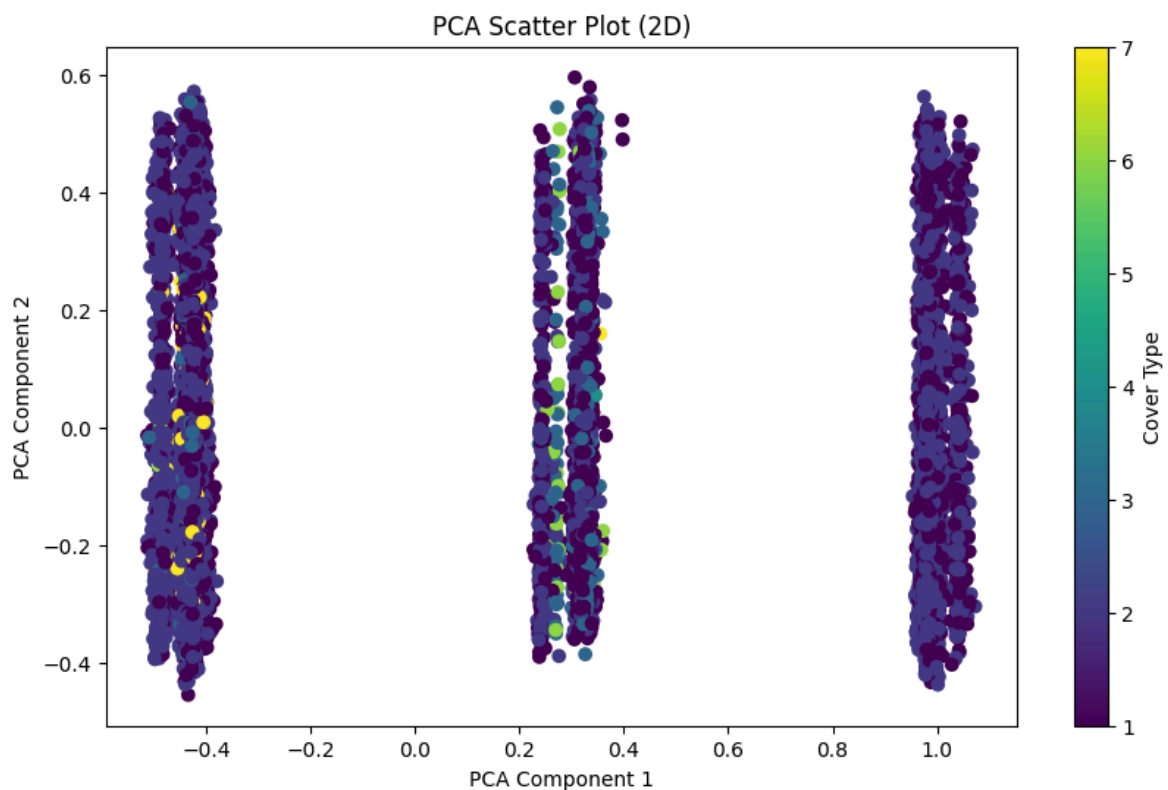
name of the column

A7 Fill the below cells. Use extra cells as per your necessary

```
In [21]: df_plot = pd.concat([data_subset_x, data_subset_y.rename('Cover_Type')], axis=1)
```


```
In [22]: # Plotting in 2D
plt.figure(figsize=(10, 6))
plt.scatter(df_plot['PCA_Component_1'], df_plot['PCA_Component_2'], c=df_plot['Cover_Type'])
plt.title('PCA Scatter Plot (2D)')
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.colorbar(label='Cover Type')
plt.show()

# Plotting in 3D
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')
scatter = ax.scatter(df_plot['PCA_Component_1'], df_plot['PCA_Component_2'], df_plot['PCA_Component_3'], c=df_plot['Cover_Type'])
ax.set_title('PCA Scatter Plot (3D)')
ax.set_xlabel('PCA Component 1')
ax.set_ylabel('PCA Component 2')
ax.set_zlabel('PCA Component 3')
plt.colorbar(scatter, label='Cover Type')
plt.show()
```



PCA Scatter Plot (3D)





Q8: Now we will plot all points from our dataframe `df_plot` Using the result from **PCA**

1. Use `pca-one` and `pca-two` column as X and Y axis respectively.
2. Use seaborn scatterplot for plotting the points.


Note: Use the notebook from class for reference. The link is provided below.

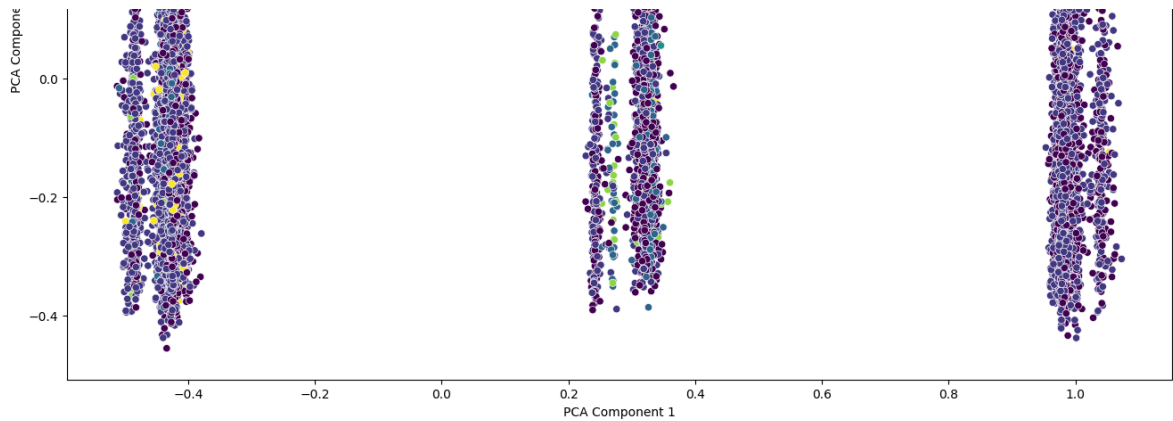
Link: https://git.txstate.edu/ML/2024Spring/tree/main/project/examples/Data_Viz_with_PC

A8 Fill the below cells. Use extra cells as per your necessary

In [23]:

```
plt.figure(figsize=(16, 10))
sns.scatterplot(data=df_plot, x='PCA_Component_1', y='PCA_Component_2', hue='Cover Type')
plt.title('Scatter Plot using PCA Components (2D)')
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.legend(title='Cover Type')
plt.show()
```





Q9: Now we will plot all points from our dataframe `df_plot` Using result from T-SNE.

1. Use `tsne-2d-one` and `tsne-2d-one` column as X and Y axis respectively.
2. Use seaborn scatterplot for plotting the points.

Note: Use the notebook from class for reference. The link is provided below.

Link:

https://git.txstate.edu/ML/2024Spring/tree/main/project/examples/Data_Viz_with_PCA_TSNE.

A9 Fill the below cells. Use extra cells as per your necessary

In [24]:

```
plt.figure(figsize=(16, 10))
sns.scatterplot(data=df_plot, x='tSNE_Component_1', y='tSNE_Component_2', hue='Cover Type')
plt.title('Scatter Plot using t-SNE Components (2D)')
plt.xlabel('tSNE Component 1')
plt.ylabel('tSNE Component 2')
plt.legend(title='Cover Type')
plt.show()
```



Part 2: Data Analysis and Classification Using Entire Dataset

Q10: Observe the data plotting and find the relation between datapoints and their characteristics.

1. Reduce the dimension of our `Scaled_X` dataframe to 3 using PCA algorithm.
2. Store the result into a variable named `pca_result`
3. Create Train data and Test data using the `pca_result` and `Y`.

Note:

1. Consider `pca_result` as X values, and `Y` as y values.
2. You can use `sklearn train_test_split`
3. Keep Train and Test ratio as : 75%:25%

A10 Fill the below cells. Use extra cells as per your necessary

```
In [25]: #You can create or remove cells as per your need
pca = PCA(n_components=3)
pca_result = pca.fit_transform(Scaled_X)
```

```
In [26]: x_train, x_test, y_train, y_test = train_test_split(pca_result, Y, test_size=0.
```

Now, Select Three best model for our dataset. You have to decide three models which might work well with our dataset.

Q11

Model Number 1

1. Reason behind choosing the model.
2. Create the model using `sklearn` or any proper library
3. Fit the model with the train data
4. Get the score from the model using test data

A11 Fill the below cells. Use extra cells as per your necessary

Answer for Q.No:1 goes here Random Forest is a versatile and powerful ensemble learning method that is known for its robustness and ability to handle high-dimensional data with complex relationships.

```
In [27]: from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import accuracy_score

# Initialize Random Forest Classifier
rf_model = RandomForestClassifier(random_state=42)

# Fit the model with the training data
rf_model.fit(x_train, y_train)

# Get the accuracy score from the model using the test data
rf_score = rf_model.score(x_test, y_test)
print("Accuracy score for Random Forest Classifier:", rf_score)
```

Accuracy score for Random Forest Classifier: 0.60825

Q12

Model Number 2

1. Reason behind choosing the model.
2. Create the model using sklearn or any proper library
3. Fit the model with the train data
4. Get the score from the model using test data

A12 Fill the below cells. Use extra cells as per your necessary Replace ??? with code in the code cell below

Answer for Q.No:1 goes here SVMs are powerful models known for their effectiveness in high-dimensional spaces, making them suitable for our dataset. They work well with complex relationships and can handle both linear and non-linear data.

In [28]:

```
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

# Initialize Support Vector Machine Classifier
svm_model = SVC()

# Fit the model with the training data
svm_model.fit(x_train, y_train)

# Get the accuracy score from the model using the test data
svm_score = svm_model.score(x_test, y_test)
print("Accuracy score for Support Vector Machine Classifier:", svm_score)
```

Accuracy score for Support Vector Machine Classifier: 0.5886

Q13

Model Number 3

1. Reason behind choosing the model.
2. Create the model using sklearn or any proper library
3. Fit the model with the train data
4. Get the score from the model using test data

A13 Fill the below cells. Use extra cells as per your necessary

Answer for Q.No:1 goes here Gradient Boosting is an ensemble learning method that builds multiple decision trees sequentially, where each tree corrects the errors of the previous one. It is known for its high predictive accuracy and ability to handle complex relationships in the data.

In [29]:

```
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.metrics import accuracy_score

# Initialize Gradient Boosting Classifier
gb_model = GradientBoostingClassifier()

# Fit the model with the training data
gb_model.fit(x_train, y_train)

# Get the accuracy score from the model using the test data
gb_score = gb_model.score(x_test, y_test)
print("Accuracy score for Gradient Boosting Classifier:", gb_score)
```

Accuracy score for Gradient Boosting Classifier: 0.6114

Q14

1. Plot a histogram using Y dataframe and display the per-class data distribution(number of rows per class).
2. Also print the number of rows per class as numeric value.


A14 Fill the below cells. Use extra cells as per your necessary

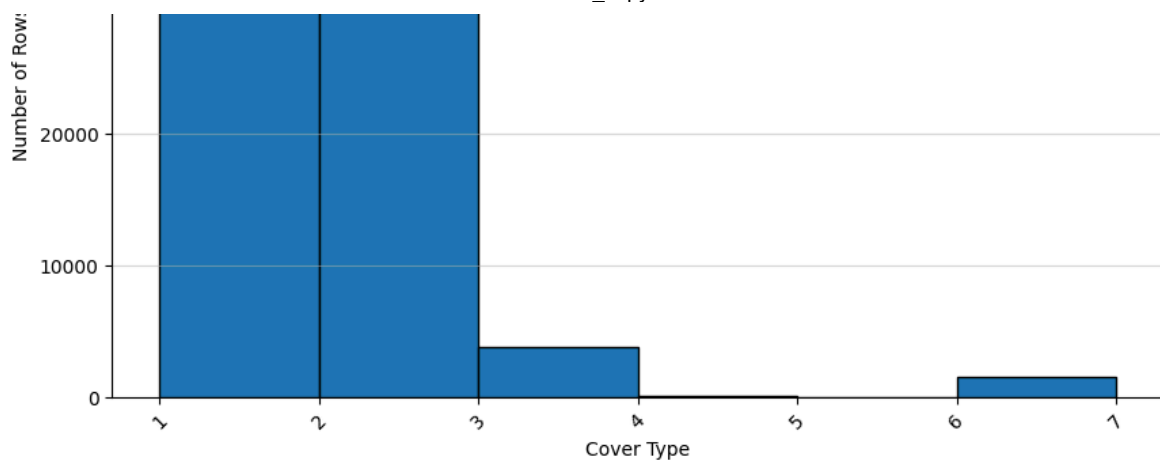
In [30]:

```
import matplotlib.pyplot as plt

# Plot histogram
plt.figure(figsize=(10, 6))
plt.hist(Y, bins=len(Y.unique()), edgecolor='black')
plt.title('Class Distribution')
plt.xlabel('Cover Type')
plt.ylabel('Number of Rows')
plt.xticks(rotation=45)
plt.grid(axis='y', alpha=0.5)
plt.show()

# Print number of rows per class
print("Number of rows per class:")
print(Y.value_counts())
```





Number of rows per class:

2 45393

1 29311

3 3816

7 1245

6 229

4 6

Name: Cover_Type, dtype: int64

Q15

1. From the histogram we can see that the dataset is highly imbalanced.
2. Use a proper dataset balancing technique to make the dataset balanced.
3. Plot a histogram using new y values and display the per-class data distribution(number of rows per class).

Note: Use can use the `imblearn.over_sampling` library for this task. But use appropriate strategy for the method.

Follow the documentation for details: https://imbalanced-learn.org/stable/references/generated/imblearn.over_sampling.SMOTE.html

A15 Fill the below cells. Use extra cells as per your necessary

```
In [31]: from imblearn.over_sampling import SMOTE

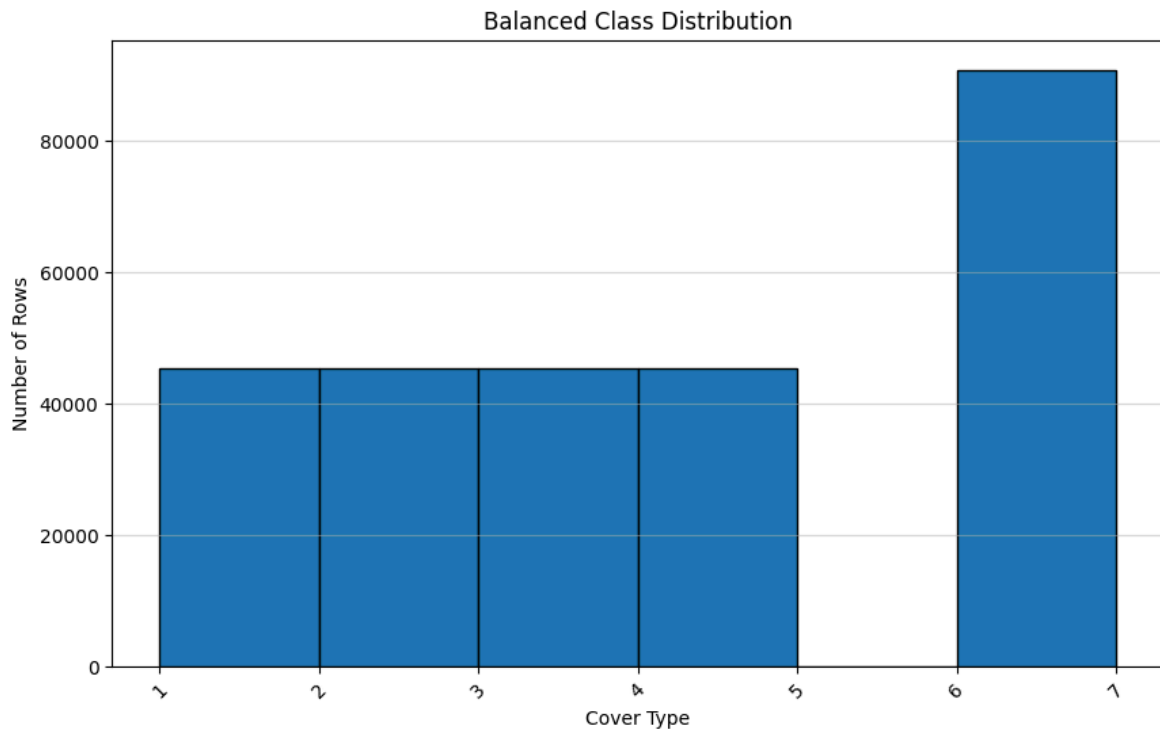
# Initialize SMOTE
smote = SMOTE(sampling_strategy='auto', random_state=42)

# Resample the dataset
X_res, y_res = smote.fit_resample(Scaled_X, Y)
```

```
In [32]: import matplotlib.pyplot as plt

# Plot histogram for balanced data
plt.figure(figsize=(10, 6))
plt.hist(y_res, bins=len(y_res.unique()), edgecolor='black')
plt.title('Balanced Class Distribution')
plt.xlabel('Cover Type')
plt.ylabel('Number of Rows')
plt.xticks(rotation=45)
```

```
plt.grid(axis='y', alpha=0.5)
plt.show()
```



Q16

1. Create new Train and Test data from the balanced X and Y value.
2. Keep Train and Test ratio as : 75%:25%

A16 Fill the below cells. Use extra cells as per your necessary

In [33]:

```
from sklearn.model_selection import train_test_split

# Split the balanced data into train and test sets
x_train, x_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.2
```

Q17

Now, Use the previously initialized three models and calculate the score from our new balanced dataset.

Model Number 1

1. Fit the model with the new train data(Use the previous Model 1)
2. Get the score from the model using new test data

A17 Fill the below cells. Use extra cells as per your necessary

In [34]:

```
# Fit the Random Forest model with the new train data
rf_model.fit(x_train, y_train)
```

```
# Get the accuracy score from the model using the new test data
rf_score_balanced = rf_model.score(x_test, y_test)
print("Accuracy score for Random Forest Classifier with balanced data:", rf_sco
```

Accuracy score for Random Forest Classifier with balanced data: 0.9824643853723014

Model Number 2

1. Fit the model with the new train data(Use the previous Model 2)
2. Get the score from the model using new test data

Fill the below cells. Use extra cells as per your necessary

In [35]:

```
# Fit the Support Vector Machine model with the new train data
svm_model.fit(x_train, y_train)

# Get the accuracy score from the model using the new test data
svm_score_balanced = svm_model.score(x_test, y_test)
print("Accuracy score for Support Vector Machine Classifier with balanced data:
```

Accuracy score for Support Vector Machine Classifier with balanced data: 0.9521956234395653

Model Number 3

1. Fit the model with the new train data(Use the previous Model 3)
2. Get the score from the model using new test data

Fill the below cells. Use extra cells as per your necessary

In [36]:

```
# Fit the Gradient Boosting model with the new train data
gb_model.fit(x_train, y_train)

# Get the accuracy score from the model using the new test data
gb_score_balanced = gb_model.score(x_test, y_test)
print("Accuracy score for Gradient Boosting Classifier with balanced data:", gb
```

Accuracy score for Gradient Boosting Classifier with balanced data: 0.9388015861359965

After making the dataset balanced we can see a significant improve in the performance for all three models.

